

analyze_training_multi_5_200

November 10, 2025

1 SUMO Traffic Light Control - Training Analysis

This notebook analyzes training metrics from DQN agent training runs. It loads metrics from the `output/` directory and provides comprehensive visualizations.

```
[1]: import json
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from pathlib import Path
from typing import Dict, List, Optional
import glob

# Set style for better-looking plots
plt.style.use('seaborn-v0_8-darkgrid')
sns.set_palette("husl")

# Configuration
OUTPUT_DIR = Path("output")

print("Libraries imported successfully!")
```

Libraries imported successfully!

1.1 Load Training Metrics

Select a run ID or load the most recent run automatically.

```
[ ]: def load_metrics(run_id: Optional[str] = None) -> tuple[Dict, Dict, Path]:
    """Load metrics and config for a specific run."""
    if run_id is None:
        # Find most recent run
        runs = sorted(OUTPUT_DIR.glob("*/*"), key=lambda p: p.stat().st_mtime,
            ↪reverse=True)
        if not runs:
            raise ValueError(f"No runs found in {OUTPUT_DIR}")
        run_dir = runs[0]
        run_id = run_dir.name
```

```

        print(f"Loading most recent run: {run_id}")
    else:
        run_dir = OUTPUT_DIR / run_id

    metrics_file = run_dir / "metrics.json"
    config_file = run_dir / "config.json"

    if not metrics_file.exists():
        raise FileNotFoundError(f"Metrics file not found: {metrics_file}")

    with open(metrics_file) as f:
        metrics = json.load(f)

    config = {}
    if config_file.exists():
        with open(config_file) as f:
            config = json.load(f)

    return metrics, config, run_dir

# Load metrics (change run_id to analyze a specific run)
# Set to None to load most recent, or specify a run_id like
↳ "20251109_204535_7ead7c5a"
RUN_ID = "20251109_234910_4da3be4b" # Change to specific run_id or None for
↳ most recent
metrics, config, run_dir = load_metrics(RUN_ID)

df = None # Will hold episode-level DataFrame once constructed

print(f"\nLoaded metrics for run: {metrics.get('run_id', 'unknown')}")
print(f"Mode: {metrics.get('mode', 'unknown')}")
print(f"Episodes: {metrics.get('episodes', 0)}")
if metrics.get('mode') == 'multi_agent_training':
    print(f"Number of agents: {metrics.get('num_agents', 0)}")
    print(f"Agent IDs: {metrics.get('tls_ids', [])}")

```

```

Loaded metrics for run: 20251109_234910_4da3be4b
Mode: multi_agent_training
Episodes: 200
Number of agents: 5
Agent IDs: ['group_0', 'group_1', 'group_2', 'group_3', 'group_4']

```

1.2 Learning Diagnostics

Check if the agent is learning or if metrics are flat (indicating no learning progress).

```

[3]: if df is not None:
    # Check for flat metrics (no learning)
    reward_col = 'reward' if 'reward' in df.columns else 'total_reward'
    reward_std = df[reward_col].std()
    queues_std = df['total_queues'].std() if 'total_queues' in df.columns else 0
    waiting_std = df['total_waiting_time'].std() if 'total_waiting_time' in df.
    ↪columns else 0

    reward_change = abs(df[reward_col].iloc[-1] - df[reward_col].iloc[0])
    queues_change = abs(df['total_queues'].iloc[-1] - df['total_queues'].
    ↪iloc[0]) if 'total_queues' in df.columns else 0
    waiting_change = abs(df['total_waiting_time'].iloc[-1] -
    ↪df['total_waiting_time'].iloc[0]) if 'total_waiting_time' in df.columns else
    ↪0

    # Thresholds for detecting flat metrics
    REWARD_THRESHOLD = 100 # Minimum change in reward to indicate learning
    QUEUES_THRESHOLD = 1 # Minimum change in queues
    WAITING_THRESHOLD = 10 # Minimum change in waiting time

    print("=" * 80)
    print("LEARNING DIAGNOSTICS")
    print("=" * 80)

    issues = []

    if reward_std < 1.0 or reward_change < REWARD_THRESHOLD:
        issues.append(" REWARD is FLAT - No improvement detected!")
        print(f"\n WARNING: Reward metrics show NO LEARNING!")
        print(f" - Reward std: {reward_std:.2f} (very low variation)")
        print(f" - Reward change: {reward_change:.2f} (should be >
    ↪{REWARD_THRESHOLD})")
        print(f" - Initial reward: {df[reward_col].iloc[0]:.2f}")
        print(f" - Final reward: {df[reward_col].iloc[-1]:.2f}")

    if queues_std < 0.1 or queues_change < QUEUES_THRESHOLD:
        issues.append(" QUEUES are FLAT - No improvement detected!")
        print(f"\n WARNING: Queue metrics show NO LEARNING!")
        print(f" - Queue std: {queues_std:.2f} (very low variation)")
        print(f" - Queue change: {queues_change:.0f} (should be >
    ↪{QUEUES_THRESHOLD})")
        print(f" - Initial queues: {df['total_queues'].iloc[0]:.0f}")
        print(f" - Final queues: {df['total_queues'].iloc[-1]:.0f}")

    if waiting_std < 1.0 or waiting_change < WAITING_THRESHOLD:
        issues.append(" WAITING TIME is FLAT - No improvement detected!")
        print(f"\n WARNING: Waiting time metrics show NO LEARNING!")

```

```

        print(f"    - Waiting std: {waiting_std:.2f} (very low variation)")
        print(f"    - Waiting change: {waiting_change:.2f} (should be >_{
↪{WAITING_THRESHOLD}})")
        print(f"    - Initial waiting: {df['total_waiting_time'].iloc[0]:.2f}s")
        print(f"    - Final waiting: {df['total_waiting_time'].iloc[-1]:.2f}s")

    if issues:
        print("\n" + "=" * 80)
        print("    CRITICAL: Agent shows NO LEARNING PROGRESS!")
        print("=" * 80)
        print("\nPossible causes:")
        print("1.    Traffic too heavy - Agent cannot learn (check flow_rate_{
↪parameter}")
        print("2.    Learning rate too low - Agent learns too slowly")
        print("3.    Reward function not working - Agent gets no useful_{
↪feedback}")
        print("4.    Network too small - Agent cannot represent good policies")
        print("5.    Exploration too low - Agent stuck in local minimum")
        print("\n Recommendations:")
        print("    • Retrain with CORRECTED flow_rate (use --insertion-rate_{
↪instead of -p)")
        print("    • Try lower flow_rate (e.g., 400-600 vehicles/hour)")
        print("    • Check if reward function is providing meaningful signals")
        print("    • Verify network architecture is appropriate for the problem")
        print("    • Consider using ImprovedDQNAgent (--agent-type improved)")
        print("=" * 80)
    else:
        print("\n Learning progress detected - metrics are changing over time")
        print(f"    - Reward improved by: {reward_change:.2f}")
        print(f"    - Queues changed by: {queues_change:.0f}")
        print(f"    - Waiting time changed by: {waiting_change:.2f}s")

    print()

```

1.3 Convert to DataFrame

Convert episode data to a pandas DataFrame for easier analysis.

```

[4]: if metrics.get('mode') in ['training', 'multi_agent_training', _
↪'centralized_training'] and 'episode_data' in metrics:
    df = pd.DataFrame(metrics['episode_data'])

    # Handle different modes
    if metrics.get('mode') == 'multi_agent_training':
        # Multi-agent mode: use total_reward and avg_epsilon
        if 'total_reward' in df.columns:
            df['reward'] = df['total_reward'] # Alias for compatibility

```

```

        if 'avg_epsilon' in df.columns:
            df['epsilon'] = df['avg_epsilon'] # Alias for compatibility
        elif metrics.get('mode') == 'centralized_training':
            # Centralized mode: same structure as single-agent
            pass
        # Single-agent mode: already has 'reward' and 'epsilon'

        # Calculate additional metrics (only for single-agent mode with per-lane
        ↪data)
        if 'queues' in df.columns and df['queues'].notna().any():
            df['avg_queue_per_lane'] = df['queues'].apply(lambda x: np.mean(x) if
            ↪isinstance(x, list) and x else 0)
            df['max_queue'] = df['queues'].apply(lambda x: max(x) if isinstance(x,
            ↪list) and x else 0)
            if 'waiting_times' in df.columns and df['waiting_times'].notna().any():
                df['avg_waiting_per_lane'] = df['waiting_times'].apply(lambda x: np.
                ↪mean(x) if isinstance(x, list) and x else 0)
                df['max_waiting'] = df['waiting_times'].apply(lambda x: max(x) if
                ↪isinstance(x, list) and x else 0)

        # Calculate moving averages
        window = min(10, len(df) // 4)
        if 'reward' in df.columns:
            df['reward_ma'] = df['reward'].rolling(window=window, center=True).
            ↪mean()
            if 'total_queues' in df.columns:
                df['queues_ma'] = df['total_queues'].rolling(window=window,
                ↪center=True).mean()
            if 'total_waiting_time' in df.columns:
                df['waiting_ma'] = df['total_waiting_time'].rolling(window=window,
                ↪center=True).mean()

        if 'avg_loss' in df.columns:
            df['loss_ma'] = df['avg_loss'].rolling(window=window, center=True).
            ↪mean()

        print(f"DataFrame shape: {df.shape}")
        print(f"\nFirst few episodes:")
        cols = ['episode']
        if 'reward' in df.columns:
            cols.append('reward')
        if 'epsilon' in df.columns:
            cols.append('epsilon')
        if 'total_queues' in df.columns:
            cols.append('total_queues')
        if 'total_waiting_time' in df.columns:

```

```

        cols.append('total_waiting_time')
    print(df[cols].head())
    print(f"\nLast few episodes:")
    print(df[cols].tail())
else:
    print("No training episode data found. This might be an evaluation-only run.
↪")
    df = None

```

DataFrame shape: (200, 16)

First few episodes:

	episode	reward	epsilon	total_queues	total_waiting_time
0	1	108049.3	0.910999	0	0.0
1	2	108815.6	0.821899	0	0.0
2	3	109644.1	0.732799	0	0.0
3	4	108028.0	0.643699	2	4.0
4	5	108627.3	0.554599	2	5.0

Last few episodes:

	episode	reward	epsilon	total_queues	total_waiting_time
195	196	-237397.9	0.01	62	5899.0
196	197	-228640.3	0.01	39	1249.0
197	198	-90149.2	0.01	30	1758.0
198	199	-74811.8	0.01	28	1213.0
199	200	-25335.1	0.01	32	2476.0

1.4 Training Overview

Display key statistics about the training run.

```

[5]: if df is not None:
    print("=" * 80)
    print("TRAINING OVERVIEW")
    print("=" * 80)
    print(f"\nRun ID: {metrics.get('run_id')}")
    print(f"Total Episodes: {len(df)}")

    reward_col = 'reward' if 'reward' in df.columns else 'total_reward'
    epsilon_col = 'epsilon' if 'epsilon' in df.columns else 'avg_epsilon'

    print(f"\nReward Statistics:")
    print(f"  Initial (first 10): {df.head(10)[reward_col].mean():.2f}")
    print(f"  Final (last 10): {df.tail(10)[reward_col].mean():.2f}")
    print(f"  Best: {df[reward_col].max():.2f}")
    print(f"  Worst: {df[reward_col].min():.2f}")
    print(f"  Mean: {df[reward_col].mean():.2f}")
    print(f"  Std: {df[reward_col].std():.2f}")

```

```

if 'total_queues' in df.columns:
    print(f"\nQueue Statistics:")
    print(f"  Initial (first 10): {df.head(10)['total_queues'].mean():.1f}")
    print(f"  Final (last 10): {df.tail(10)['total_queues'].mean():.1f}")
    print(f"  Best (lowest): {df['total_queues'].min():.0f}")
    print(f"  Worst (highest): {df['total_queues'].max():.0f}")

if 'total_waiting_time' in df.columns:
    print(f"\nWaiting Time Statistics:")
    print(f"  Initial (first 10): {df.head(10)['total_waiting_time'].mean():.2f}s")
    print(f"  Final (last 10): {df.tail(10)['total_waiting_time'].mean():.2f}s")
    print(f"  Best (lowest): {df['total_waiting_time'].min():.2f}s")

if epsilon_col in df.columns:
    print(f"\nEpsilon (Exploration):")
    print(f"  Start: {df[epsilon_col].iloc[0]:.3f}")
    print(f"  End: {df[epsilon_col].iloc[-1]:.3f}")

# Multi-agent specific stats
if metrics.get('mode') == 'multi_agent_training' and 'agent_rewards' in df.columns:
    print(f"\nMulti-Agent Statistics:")
    print(f"  Number of agents: {metrics.get('num_agents', 0)}")
    if 'avg_reward_per_agent' in df.columns:
        print(f"  Avg reward per agent (final): {df.tail(10)['avg_reward_per_agent'].mean():.2f}")

if 'avg_loss' in df.columns and df['avg_loss'].notna().any():
    loss_data = df[df['avg_loss'].notna()]
    if len(loss_data) > 0:
        print(f"\nLoss Statistics:")
        print(f"  Mean: {loss_data['avg_loss'].mean():.6f}")
        print(f"  Final: {loss_data['avg_loss'].iloc[-1]:.6f}")

if 'best_eval_reward' in metrics and metrics['best_eval_reward']:
    print(f"\nBest Evaluation Reward: {metrics['best_eval_reward']:.2f}")

```

TRAINING OVERVIEW

Run ID: 20251109_234910_4da3be4b
Total Episodes: 200

Reward Statistics:

Initial (first 10): 102515.22
Final (last 10): -137917.70
Best: 109644.10
Worst: -287687.20
Mean: 17837.04
Std: 82885.55

Queue Statistics:

Initial (first 10): 2.3
Final (last 10): 43.7
Best (lowest): 0
Worst (highest): 80

Waiting Time Statistics:

Initial (first 10): 34.30s
Final (last 10): 3929.00s
Best (lowest): 0.00s

Epsilon (Exploration):

Start: 0.911
End: 0.010

Multi-Agent Statistics:

Number of agents: 5
Avg reward per agent (final): -27583.54

1.5 Visualization: Reward Over Time

Plot episode rewards to see learning progress.

```
[6]: if df is not None:
    fig, ax = plt.subplots(figsize=(12, 6))

    reward_col = 'reward' if 'reward' in df.columns else 'total_reward'
    title_suffix = ' (Multi-Agent)' if metrics.get('mode') == '
    ↪ 'multi_agent_training' else ''

    ax.plot(df['episode'], df[reward_col], alpha=0.3, label='Episode Reward',
    ↪ color='blue', linewidth=0.5)
    if 'reward_ma' in df.columns:
        ax.plot(df['episode'], df['reward_ma'], label=f'Moving Average
        ↪ (window={window})',
                color='red', linewidth=2)

    ax.set_xlabel('Episode', fontsize=12)
    ax.set_ylabel('Reward', fontsize=12)
```




```

ax.set_title(f'Training Reward Over Time{title_suffix}', fontsize=14,
fontweight='bold')
ax.legend()
ax.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

```



1.6 Visualization: Queue Length Over Time

Track how queue lengths change during training.

```


[7]: if df is not None:
    fig, ax = plt.subplots(figsize=(12, 6))

    ax.plot(df['episode'], df['total_queues'], alpha=0.3, label='Total Queues',
            color='orange', linewidth=0.5)
    if 'queues_ma' in df.columns:
        ax.plot(df['episode'], df['queues_ma'], label=f'Moving Average (window={window})',
                color='red', linewidth=2)

    ax.set_xlabel('Episode', fontsize=12)
    ax.set_ylabel('Total Queue Length', fontsize=12)
    ax.set_title('Total Queue Length Over Time', fontsize=14, fontweight='bold')
    ax.legend()
    ax.grid(True, alpha=0.3)

```

```
plt.tight_layout()
plt.show()
```



1.7 Visualization: Waiting Time Over Time


Monitor waiting times to assess traffic flow improvement.

```
[8]: if df is not None:
    fig, ax = plt.subplots(figsize=(12, 6))

    ax.plot(df['episode'], df['total_waiting_time'], alpha=0.3, label='Total_
    ↳Waiting Time',
            color='green', linewidth=0.5)
    if 'waiting_ma' in df.columns:
        ax.plot(df['episode'], df['waiting_ma'], label=f'Moving Average_
        ↳(window={window})',
                color='red', linewidth=2)

    ax.set_xlabel('Episode', fontsize=12)
    ax.set_ylabel('Total Waiting Time (seconds)', fontsize=12)
    ax.set_title('Total Waiting Time Over Time', fontsize=14, fontweight='bold')
    ax.legend()
    ax.grid(True, alpha=0.3)

    plt.tight_layout()
    plt.show()
```



1.8 Visualization: Epsilon (Exploration Rate)


Track how exploration decreases over time as the agent becomes more confident.

```
[9]: if df is not None:
    epsilon_col = 'epsilon' if 'epsilon' in df.columns else 'avg_epsilon'
    if epsilon_col in df.columns:
        fig, ax = plt.subplots(figsize=(12, 6))

        title_suffix = ' (Multi-Agent Avg)' if metrics.get('mode') == '
        ↪'multi_agent_training' else ''
        ax.plot(df['episode'], df[epsilon_col], label='Epsilon (Exploration
        ↪Rate)',
                color='purple', linewidth=2)

        ax.set_xlabel('Episode', fontsize=12)
        ax.set_ylabel('Epsilon', fontsize=12)
        ax.set_title(f'Exploration Rate Decay{title_suffix}', fontsize=14,
        ↪fontweight='bold')
        ax.legend()
        ax.grid(True, alpha=0.3)

        plt.tight_layout()
        plt.show()
    else:
        print("Epsilon data not available in this run.")
```



1.9 Visualization: Training Loss

If available, plot the training loss to monitor learning stability.

```
[10]: if df is not None and 'avg_loss' in df.columns and df['avg_loss'].notna().any():
    loss_data = df[df['avg_loss'].notna()]

    if len(loss_data) > 0:
        fig, ax = plt.subplots(figsize=(12, 6))

        ax.plot(loss_data['episode'], loss_data['avg_loss'], alpha=0.3,
                label='Episode Loss', color='darkred', linewidth=0.5)
        if 'loss_ma' in loss_data.columns:
            ax.plot(loss_data['episode'], loss_data['loss_ma'],
                    label=f'Moving Average (window={window})', color='red',
                    linewidth=2)

        ax.set_xlabel('Episode', fontsize=12)
        ax.set_ylabel('Loss', fontsize=12)
        ax.set_title('Training Loss Over Time', fontsize=14, fontweight='bold')
        ax.legend()
        ax.grid(True, alpha=0.3)

        plt.tight_layout()
        plt.show()
    else:
        print("No loss data available.")
else:
    print("Loss tracking not available in this run.")
```

Loss tracking not available in this run.

1.10 Visualization: Action Distribution

Analyze which actions the agent prefers over time.

```
[11]: if df is not None and 'action_distribution' in df.columns:
    # Extract action distributions
    action_data = []
    for idx, row in df.iterrows():
        dist = row['action_distribution']
        if dist:
            for action, count in dist.items():
                action_data.append({
                    'episode': row['episode'],
                    'action': int(action),
                    'count': count,
                    'percentage': count / sum(dist.values()) * 100
                })

    if action_data:
        action_df = pd.DataFrame(action_data)

        fig, axes = plt.subplots(1, 2, figsize=(15, 6))

        # Plot 1: Action usage over time
        for action in sorted(action_df['action'].unique()):
            action_subset = action_df[action_df['action'] == action]
            axes[0].plot(action_subset['episode'], action_subset['percentage'],
                        label=f'Action {action}', linewidth=2, alpha=0.7)

        axes[0].set_xlabel('Episode', fontsize=12)
        axes[0].set_ylabel('Action Usage (%)', fontsize=12)
        axes[0].set_title('Action Distribution Over Time', fontsize=14,
        ↪fontweight='bold')
        axes[0].legend()
        axes[0].grid(True, alpha=0.3)

        # Plot 2: Overall action distribution
        overall_dist = action_df.groupby('action')['count'].sum()
        axes[1].bar(overall_dist.index, overall_dist.values, alpha=0.7)
        axes[1].set_xlabel('Action', fontsize=12)
        axes[1].set_ylabel('Total Count', fontsize=12)
        axes[1].set_title('Overall Action Distribution', fontsize=14,
        ↪fontweight='bold')
        axes[1].grid(True, alpha=0.3, axis='y')

        plt.tight_layout()
```

```

        plt.show()
    else:
        print("No action distribution data available.")
else:
    print("Action distribution tracking not available in this run.")

```

Action distribution tracking not available in this run.

```

[12]: if df is not None:
        fig, axes = plt.subplots(2, 2, figsize=(16, 12))

        reward_col = 'reward' if 'reward' in df.columns else 'total_reward'
        epsilon_col = 'epsilon' if 'epsilon' in df.columns else 'avg_epsilon'

        # Reward
        axes[0, 0].plot(df['episode'], df[reward_col], alpha=0.3, color='blue',
        ↳linewidth=0.5)
        if 'reward_ma' in df.columns:
            axes[0, 0].plot(df['episode'], df['reward_ma'], color='red',
        ↳linewidth=2)
        axes[0, 0].set_title('Reward Over Time', fontweight='bold')
        axes[0, 0].set_xlabel('Episode')
        axes[0, 0].set_ylabel('Reward')
        axes[0, 0].grid(True, alpha=0.3)

        # Queues
        if 'total_queues' in df.columns:
            axes[0, 1].plot(df['episode'], df['total_queues'], alpha=0.3,
        ↳color='orange', linewidth=0.5)
            if 'queues_ma' in df.columns:
                axes[0, 1].plot(df['episode'], df['queues_ma'], color='red',
        ↳linewidth=2)
            axes[0, 1].set_title('Total Queue Length', fontweight='bold')
            axes[0, 1].set_xlabel('Episode')
            axes[0, 1].set_ylabel('Queues')
            axes[0, 1].grid(True, alpha=0.3)
        else:
            axes[0, 1].text(0.5, 0.5, 'No queue data', ha='center', va='center',
        ↳transform=axes[0, 1].transAxes)
            axes[0, 1].set_title('Total Queue Length', fontweight='bold')

        # Waiting Time
        if 'total_waiting_time' in df.columns:
            axes[1, 0].plot(df['episode'], df['total_waiting_time'], alpha=0.3,
        ↳color='green', linewidth=0.5)
            if 'waiting_ma' in df.columns:

```

```

        axes[1, 0].plot(df['episode'], df['waiting_ma'], color='red',
↪linewidth=2)
        axes[1, 0].set_title('Total Waiting Time', fontweight='bold')
        axes[1, 0].set_xlabel('Episode')
        axes[1, 0].set_ylabel('Waiting Time (s)')
        axes[1, 0].grid(True, alpha=0.3)
    else:
        axes[1, 0].text(0.5, 0.5, 'No waiting time data', ha='center',
↪va='center', transform=axes[1, 0].transAxes)
        axes[1, 0].set_title('Total Waiting Time', fontweight='bold')

    # Epsilon
    if epsilon_col in df.columns:
        axes[1, 1].plot(df['episode'], df[epsilon_col], color='purple',
↪linewidth=2)
        axes[1, 1].set_title('Exploration Rate (Epsilon)', fontweight='bold')
        axes[1, 1].set_xlabel('Episode')
        axes[1, 1].set_ylabel('Epsilon')
        axes[1, 1].grid(True, alpha=0.3)
    else:
        axes[1, 1].text(0.5, 0.5, 'No epsilon data', ha='center', va='center',
↪transform=axes[1, 1].transAxes)
        axes[1, 1].set_title('Exploration Rate (Epsilon)', fontweight='bold')

    mode_str = f" - {metrics.get('mode', 'unknown').replace('_', ' ').title()}"
    plt.suptitle(f'Training Dashboard - Run: {metrics.get("run_id",
↪"unknown")}-{mode_str}',
                fontsize=16, fontweight='bold', y=0.995)
    plt.tight_layout()
    plt.show()

```



1.11 Per-Lane Analysis

Analyze performance for individual lanes.

```
[13]: if df is not None and 'queues' in df.columns:
    # Extract per-lane data
    num_lanes = len(df['queues'].iloc[0]) if len(df) > 0 and df['queues'].
    ↪iloc[0] else 0

    if num_lanes > 0:
        fig, axes = plt.subplots(1, 2, figsize=(16, 6))

        # Per-lane queues over time
        for lane_idx in range(num_lanes):
            lane_queues = df['queues'].apply(lambda x: x[lane_idx] if len(x) >
            ↪lane_idx else 0)
            axes[0].plot(df['episode'], lane_queues, alpha=0.6, label=f'Lane_
            ↪{lane_idx}', linewidth=1.5)

        axes[0].set_xlabel('Episode', fontsize=12)
```



```

        axes[0].set_ylabel('Queue Length', fontsize=12)
        axes[0].set_title('Queue Length by Lane Over Time', fontsize=14,
↪fontweight='bold')
        axes[0].legend()
        axes[0].grid(True, alpha=0.3)

        # Per-lane waiting times over time
        for lane_idx in range(num_lanes):
            lane_waiting = df['waiting_times'].apply(lambda x: x[lane_idx] if
↪len(x) > lane_idx else 0)
            axes[1].plot(df['episode'], lane_waiting, alpha=0.6, label=f'Lane
↪{lane_idx}', linewidth=1.5)

        axes[1].set_xlabel('Episode', fontsize=12)
        axes[1].set_ylabel('Waiting Time (seconds)', fontsize=12)
        axes[1].set_title('Waiting Time by Lane Over Time', fontsize=14,
↪fontweight='bold')
        axes[1].legend()
        axes[1].grid(True, alpha=0.3)

        plt.tight_layout()
        plt.show()
    else:
        print("No per-lane data available.")
else:
    print("Per-lane analysis not available.")

```

Per-lane analysis not available.

1.12 Evaluation Metrics (if available)

If periodic evaluations were performed, analyze those results.

```

[14]: if 'evaluation_data' in metrics and metrics['evaluation_data']:
        eval_df = pd.DataFrame([
            {
                'episode': eval_data['episode'],
                'avg_reward': eval_data['metrics']['avg_reward'],
                'avg_queues': eval_data['metrics']['avg_total_queues'],
                'avg_waiting': eval_data['metrics']['avg_total_waiting_time']
            }
        ])

        fig, axes = plt.subplots(1, 3, figsize=(18, 5))

        axes[0].plot(eval_df['episode'], eval_df['avg_reward'], marker='o',
↪linewidth=2, markersize=8)

```

```

axes[0].set_title('Evaluation Reward', fontweight='bold')
axes[0].set_xlabel('Episode')
axes[0].set_ylabel('Average Reward')
axes[0].grid(True, alpha=0.3)

axes[1].plot(eval_df['episode'], eval_df['avg_queues'], marker='o',
↳linewidth=2, markersize=8, color='orange')
axes[1].set_title('Evaluation Queues', fontweight='bold')
axes[1].set_xlabel('Episode')
axes[1].set_ylabel('Average Queues')
axes[1].grid(True, alpha=0.3)

axes[2].plot(eval_df['episode'], eval_df['avg_waiting'], marker='o',
↳linewidth=2, markersize=8, color='green')
axes[2].set_title('Evaluation Waiting Time', fontweight='bold')
axes[2].set_xlabel('Episode')
axes[2].set_ylabel('Average Waiting Time (s)')
axes[2].grid(True, alpha=0.3)

plt.suptitle('Periodic Evaluation Results', fontsize=14, fontweight='bold')
plt.tight_layout()
plt.show()

print("\nEvaluation Summary:")
print(eval_df.describe())
else:
    print("No evaluation data available (--eval-every was not used).")

```

No evaluation data available (--eval-every was not used).

1.13 Configuration Summary

Display the training configuration used for this run.

```

[15]: print("=" * 80)
print("TRAINING CONFIGURATION")
print("=" * 80)

if config:
    for key, value in config.items():
        print(f"{key:25s}: {value}")
else:
    print("Configuration file not found.")

if 'config' in metrics:
    print("\nAdditional Metrics Config:")
    for key, value in metrics['config'].items():
        print(f"{key:25s}: {value}")

```

```
=====
TRAINING CONFIGURATION
=====
```

```
run_id           : 20251109_234910_4da3be4b
episodes         : 200
max_steps        : 900
step_length      : 1.0
seed             : None
learning_rate    : 0.001
batch_size       : 64
replay_size      : 50000
target_update    : 200
eval_every       : 0
eval_episodes    : 5
osm_file         : maps/sanFrancisco.osm
flow_rate        : 1000
tls_id           : None
device           : cpu
```

1.14 Compare Multiple Runs

Load and compare multiple training runs side by side.

1.15 Multi-Agent Per-Agent Analysis

For multi-agent training, visualize individual agent performance.

```
[16]: if df is not None and metrics.get('mode') == 'multi_agent_training' and
      ↪ 'agent_rewards' in df.columns:
      # Extract per-agent rewards
      agent_rewards_data = []
      for idx, row in df.iterrows():
          agent_rewards = row['agent_rewards']
          if isinstance(agent_rewards, dict):
              for agent_id, reward in agent_rewards.items():
                  agent_rewards_data.append({
                      'episode': row['episode'],
                      'agent_id': agent_id,
                      'reward': reward
                  })

      if agent_rewards_data:
          agent_rewards_df = pd.DataFrame(agent_rewards_data)

          fig, axes = plt.subplots(2, 1, figsize=(14, 10))

          # Plot 1: Per-agent rewards over time
          for agent_id in sorted(agent_rewards_df['agent_id'].unique()):
```

```

        agent_data = agent_rewards_df[agent_rewards_df['agent_id'] ==
↪agent_id]
        axes[0].plot(agent_data['episode'], agent_data['reward'],
                      label=f'{agent_id}', linewidth=2, alpha=0.7)

        axes[0].set_xlabel('Episode', fontsize=12)
        axes[0].set_ylabel('Reward', fontsize=12)
        axes[0].set_title('Per-Agent Rewards Over Time', fontsize=14,
↪fontweight='bold')
        axes[0].legend()
        axes[0].grid(True, alpha=0.3)

# Plot 2: Per-agent epsilon over time
        if 'agent_epsilon' in df.columns:
            agent_epsilon_data = []
            for idx, row in df.iterrows():
                agent_epsilon = row['agent_epsilon']
                if isinstance(agent_epsilon, dict):
                    for agent_id, epsilon in agent_epsilon.items():
                        agent_epsilon_data.append({
                            'episode': row['episode'],
                            'agent_id': agent_id,
                            'epsilon': epsilon
                        })

            if agent_epsilon_data:
                agent_epsilon_df = pd.DataFrame(agent_epsilon_data)
                for agent_id in sorted(agent_epsilon_df['agent_id'].unique()):
                    agent_data = agent_epsilon_df[agent_epsilon_df['agent_id']
↪== agent_id]

                    axes[1].plot(agent_data['episode'], agent_data['epsilon'],
                                label=f'{agent_id}', linewidth=2, alpha=0.7)

                    axes[1].set_xlabel('Episode', fontsize=12)
                    axes[1].set_ylabel('Epsilon', fontsize=12)
                    axes[1].set_title('Per-Agent Epsilon (Exploration Rate) Over Time',
↪fontsize=14, fontweight='bold')
                    axes[1].legend()
                    axes[1].grid(True, alpha=0.3)

            plt.tight_layout()
            plt.show()


# Print per-agent statistics
        print("\nPer-Agent Final Statistics (last 10 episodes):")
        print("=" * 80)
        for agent_id in sorted(agent_rewards_df['agent_id'].unique()):

```

```

        agent_data = agent_rewards_df[agent_rewards_df['agent_id'] ==
↪agent_id]
        final_episodes = agent_data.nlargest(10, 'episode')
        print(f"\n{agent_id}:")
        print(f"  Final avg reward: {final_episodes['reward'].mean():.2f}")
        print(f"  Best reward: {agent_data['reward'].max():.2f}")
        print(f"  Worst reward: {agent_data['reward'].min():.2f}")
    else:
        print("No per-agent reward data available.")
elif df is not None and metrics.get('mode') != 'multi_agent_training':
    print("This is not a multi-agent training run. Per-agent analysis not
↪applicable.")
else:
    print("No data available for per-agent analysis.")

```



Per-Agent Final Statistics (last 10 episodes):

```

=====
group_0:
  Final avg reward: -34241.11
  Best reward: 24580.80

```

Worst reward: -187807.40

group_1:

Final avg reward: -30761.82

Best reward: 25979.10

Worst reward: -201139.20

group_2:

Final avg reward: -77325.74

Best reward: 25438.70

Worst reward: -161123.90

group_3:

Final avg reward: -4332.41

Best reward: 25607.60

Worst reward: -106297.60

group_4:

Final avg reward: 8743.38

Best reward: 8934.60

Worst reward: -38646.60

```
[17]: def compare_runs(run_ids: List[str]) -> None:
        """Compare multiple training runs."""
        fig, axes = plt.subplots(2, 2, figsize=(16, 12))

        for run_id in run_ids:
            try:
                metrics, config, _ = load_metrics(run_id)
                if metrics.get('mode') == 'training' and 'episode_data' in metrics:
                    df = pd.DataFrame(metrics['episode_data'])

                    window = min(10, len(df) // 4)
                    df['reward_ma'] = df['reward'].rolling(window=window,
↪center=True).mean()
                    df['queues_ma'] = df['total_queues'].rolling(window=window,
↪center=True).mean()

                    label = f"{run_id[:15]}..."

                    axes[0, 0].plot(df['episode'], df['reward_ma'], label=label,
↪linewidth=2)
                    axes[0, 1].plot(df['episode'], df['queues_ma'], label=label,
↪linewidth=2)
                    axes[1, 0].plot(df['episode'], df['total_waiting_time'],
↪label=label, linewidth=2, alpha=0.7)
```

```

        axes[1, 1].plot(df['episode'], df['epsilon'], label=label,
↳linewidth=2)
    except Exception as e:
        print(f"Error loading run {run_id}: {e}")

    axes[0, 0].set_title('Reward Comparison', fontweight='bold')
    axes[0, 0].set_xlabel('Episode')
    axes[0, 0].set_ylabel('Reward (MA)')
    axes[0, 0].legend()
    axes[0, 0].grid(True, alpha=0.3)

    axes[0, 1].set_title('Queue Comparison', fontweight='bold')
    axes[0, 1].set_xlabel('Episode')
    axes[0, 1].set_ylabel('Queues (MA)')
    axes[0, 1].legend()
    axes[0, 1].grid(True, alpha=0.3)

    axes[1, 0].set_title('Waiting Time Comparison', fontweight='bold')
    axes[1, 0].set_xlabel('Episode')
    axes[1, 0].set_ylabel('Waiting Time (s)')
    axes[1, 0].legend()
    axes[1, 0].grid(True, alpha=0.3)

    axes[1, 1].set_title('Epsilon Comparison', fontweight='bold')
    axes[1, 1].set_xlabel('Episode')
    axes[1, 1].set_ylabel('Epsilon')
    axes[1, 1].legend()
    axes[1, 1].grid(True, alpha=0.3)

    plt.tight_layout()
    plt.show()

# Uncomment and modify to compare specific runs:
# compare_runs(["run_id_1", "run_id_2", "run_id_3"])

```